



Operating Systems

Deadlock




Deadlock

A set of processes is in a deadlocked state if every process in the set is **waiting for an event** that can be caused only by another process in the set.

Example:

- Two process : P1, P2
- Two resource: R1, R2
- P1 is holding R1 and waiting for R2
- P2 is holding R2 and waiting for R1



System Model

- A system contains finite number of resources
 - These resources are distributed among competing processes
 - Examples: CPU cycle, File, I/O devices (printer, DVD drive) etc.
- Each resource type may have identical number of instances
- A process may request as many resource as it requires.
 - But, it should not exceed the total number of available resources of the system
- A process may utilize a resource only in the following sequence
 1. **Request:** requests the resource
 2. **Use:** operate on the resource
 3. **Release:** releases the resource

Necessary Conditions for Deadlock

A deadlock situation can arise if the following **four conditions hold simultaneously** in a system -




1. Mutual exclusion
2. Hold and wait
3. No preemption
4. Circular wait

If one of them is not present in a system, no deadlock will arise

No deadlock, if one of them is absent.


Mutual Exclusion

At least one resource must be held in a nonsharable mode; that is, only one process at a time can use the resource



Hold and Wait

A process must be **holding at least one resource**, and **waiting** to acquire additional resources that are currently being **held** by other processes.



No Preemption

a resource can be **released** only voluntarily by the process **holding it**, after that process has **completed its task**.




Circular Wait

A set $\{P_0, P_1, \dots, P_n\}$ of waiting processes must exist such that

P_0 is waiting for a resource held by P_1 ,

P_1 is waiting for a resource held by P_2 , ..., P_{n-1} is waiting for a resource held by P_n , and

P_n is waiting for a resource held by P_0






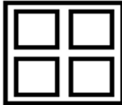


Resource Allocation Graph




Resource Allocation Graph

A set of vertices V and a set of edges E .

- V is partitioned into **two types**:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of **all the processes** in the system.
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of **all resource** types in the system.
- Request edge: directed edge from a process to a resource
 - $P_1 \rightarrow R_2$
- Assignment edge: directed edge from a resource to a process
 - $R_1 \rightarrow P_2$

Process	
Resource Type with 4 instances	
P_i requests instance of R_j	
P_i is holding an instance of R_j	

Example of Resource Allocation Graph




$P = \{P_1, P_2, P_3\}$


$R = \{R_1, R_2, R_3, R_4\}$

$E = \{P_1 \rightarrow R_1, P_2 \rightarrow R_3, R_1 \rightarrow P_2, R_2 \rightarrow P_2, R_2 \rightarrow P_1, R_3 \rightarrow P_3\}$

Resource Allocation Graph (Deadlock)



Cycle:
Possibility of Deadlock.
No Cycle:
No Deadlock.



Cycle 1: $P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$

Cycle: $P_1 \rightarrow R_1 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$

Cycle 2: $P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_2$

No Deadlock!

Deadlock!

Methods for Handling Deadlocks

- **Prevention:**
 - Ensure that the system will never enter a deadlock state
- **Avoidance:**
 - request for any resource will be granted if the resulting state of the system doesn't cause deadlock
- **Detection and recovery:**
 - Allow the system to enter a deadlock state and then recover



Operating Systems

Deadlock Prevention




Deadlock Prevention

Ensure that at least one of the necessary conditions for deadlock cannot hold in the system.

Mutual Exclusion:

- Mutual exclusion condition must hold
- At least one resource must be non-sharable
- Sharable resources do not require mutually exclusive access (example: Read only files)



Hold and Wait:

Guarantee that when a process requests a resource, it does not hold any other resources.

- Protocol 1: Request and allocate all its resources before execution.
- Protocol 2: Allow process to request resources only when the process has none.


Disadvantages:

- Low utilization of resources
- Starvation

Deadlock Prevention

No Preemption:

- If a process is **holding some resources** and **requesting** for another resource, which is held by another process, then this process needs to wait, and the resources held by it will be preempted or released.
- Preempted resources are **added to the list** of resources for which the process is waiting
- Process will be **restarted only** when it can regain its old resources, as well as the new ones that it is requesting.



Circular Wait:

- Impose **a total ordering** of all resource types
- Each process requests resources in an **increasing order** of enumeration

$R = \{R_1, R_2, \dots, R_m\}$ and we **assign unique number** to each of the resource type.

For example, $F(R_1) = 1, F(R_2) = 3, F(R_3) = 5$.

- A process can initially request any number of instances of a resource type, R_i
- After that, the process can request instances of resource type R_j if and only if $F(R_j) > F(R_i)$.
- A process requesting an instance of resource type R_j must have released any resources R_i such that $F(R_i) \geq F(R_j)$



Operating Systems

Deadlock Avoidance



Deadlock Avoidance


- System has a **priori information**, such as, maximum number of resources of each type that every **process need**
- Based on these information **construct an algorithm** that ensures that the **system** will never enter a deadlocked state
- **Dynamically** examines the **resource-allocation state** to ensure that a **circular-wait condition** can **never exist**
- **Resource-allocation state**: number of available and allocated resources, and the maximum demands of the processes.

Safe State

- System is in safe state if can allocate resources to each process (up to its maximum) in some order and still avoid a deadlock.
- Safe state = there exists a safe sequence
- Sequence $\langle P_1, P_2, \dots, P_n \rangle$ is safe if, for each P_i , the resources that P_i can still request, can be satisfied by (currently available resources + resources held by all the P_j , with $j < i$)
 - If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished.
 - When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate.
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on.

Safe State, Unsafe State & Deadlock

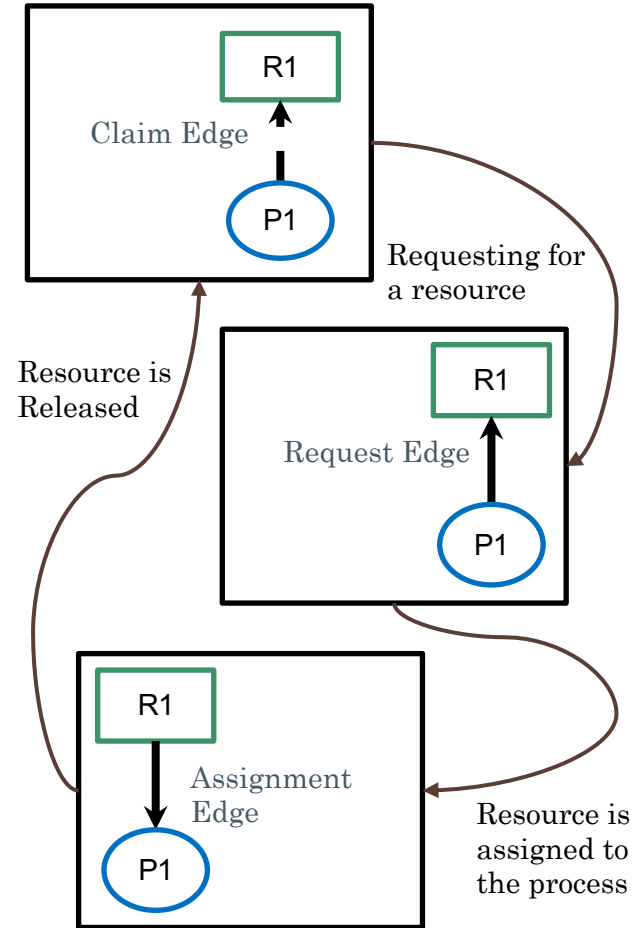
- Safe State => No Deadlock!
- Unsafe State => Possibility of Deadlock!
- To avoid deadlock, always ensure that system is in safe state.



Resource-Allocation-Graph Algorithm

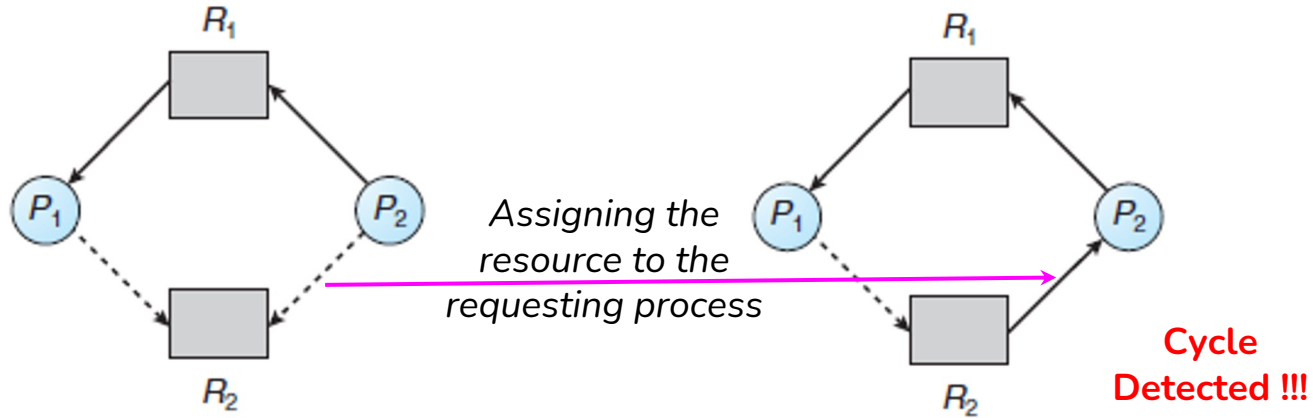
(Deadlock Avoidance Algorithm)

- Used when the system has only one instance of each resource type.
- **Request Edge** and **Assignment Edge** exist here similar to Resource allocation graph
- **Claim Edge**: claim edge $P_i \rightarrow R_j$ indicates that process P_i may request resource R_j at some time in the future
- When process P_i requests resource R_j , the claim edge $P_i \rightarrow R_j$ is converted to a request edge.
- When a resource R_j is released by P_i , the assignment edge $R_j \rightarrow P_i$ is reconverted to a claim edge $P_i \rightarrow R_j$.
- Resources must be claimed a priori in the system. i.e. before process P_i starts executing, all its claim edges must appear in the graph



Resource-Allocation-Graph Algorithm

(Deadlock Avoidance Algorithm)



Resource-allocation graph
for deadlock avoidance

An unsafe state in a
resource-allocation graph



Operating Systems

Bankers Algorithm



Banker's Algorithm (Deadlock Avoidance Algorithm)

- Multiple instances of each resource type.
- Each process must a priori claim maximum use.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.

Data Structures for the Banker's Algorithm

Need the following data structures, where **n** is the number of processes in the system and **m** is the number of resource types.

- ❑ **Available:** Array of length **m**, indicates the number of available resources of each type.
- ❑ **Max:** **n** by **m** matrix, defines **maximum demand** of each process for each type.
- ❑ **Allocation:** **n** by **m** matrix, defines number of **assigned resources** to each process for each type.
- ❑ **Need:** **n** by **m** matrix, defines number of **remaining** resources of each type needed by each process.

$$\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j]$$

Banker's Safety Algorithm

Algorithm for finding out whether or not a system is in a safe state

1. Let **Work** and **Finish** be vectors of length m and n , respectively. Initialize
Work = **Available** and
Finish[i] = **false** for $i = 0, 1, \dots, n - 1$.
2. Find an index i such that **both** of the following **meets**,
Finish[i] == **false**
Need _{i} ≤ **Work**
if no such i exists, go to step 4.
3. **Work** = **Work** + **Allocation** _{i}
Finish[i] = **true**
Go to step 2.
4. If **Finish**[i] == **true** for all i , then the system is in a safe state.

Banker's Algorithm Example

- 5 processes P0 through P4;
- 3 resource types A (10 instances), B (5 instances), and C (7 instances)
- at time T0, the following snapshot of the system:

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>		<u>Need</u>
	A B C	A B C	A B C		A B C
P ₀	0 1 0	7 5 3	3 3 2	P ₀	7 4 3
P ₁	2 0 0	3 2 2		P ₁	1 2 2
P ₂	3 0 2	9 0 2		P ₂	6 0 0
P ₃	2 1 1	2 2 2		P ₃	0 1 1
P ₄	0 0 2	4 3 3		P ₄	4 3 1

finish:

--	--	--	--	--

work:

--

```

if finish[i] == false & Needi ≤ Work then
    Work = Work + Allocationi
    Finish[i] = true
else
    wait
    
```

The system is in a safe state since the sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety criteria

Practice

	Allocation			Max			Available		
	A	B	C	A	B	C	A	B	C
P ₀	1	0	1	2	1	1	2	1	1
P ₁	2	1	2	5	4	4			
P ₂	3	0	0	3	1	1			
P ₃	1	0	1	1	1	1			

Resource-Request Algorithm for Process P_i

Request = request vector for process P_i

If **Request**[i][j] = k then process P_i wants k instances of resource type R_j

1. If $\text{Request}_i \leq \text{Need}_i$ go to step 2.
Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If $\text{Request}_i \leq \text{Available}$, go to step 3.
Otherwise P_i must wait, since resources are not available.
3. Pretend to allocate requested resources to P_i by modifying the state as follows:
 $\text{Available} = \text{Available} - \text{Request}_i;$
 $\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i;$
 $\text{Need}_i = \text{Need}_i - \text{Request}_i;$
4. Check **Bankers Safety** Algorithm, if this new state is safe.
 - If safe the resources are allocated to P_i .
 - If unsafe P_i must wait, and the old resource-allocation state is restored

What if P_1 Request (1,0,2) more ?

- To decide whether this request can be immediately granted,
 - we first check that $\mathbf{Request}_1 \leq \mathbf{Need}_1$ — that is, $(1,0,2) \leq (1,2,2)? \Rightarrow \text{TRUE}$
 - Then we check that $\mathbf{Request}_1 \leq \mathbf{Available}$ — that is, $(1,0,2) \leq (3,3,2)? \Rightarrow \text{TRUE}$
 - If both is true,

$$\mathbf{Available} = (3,3,2) - (1,0,2) = (2,3,0)$$

$$\mathbf{Allocation}_1 = (2,0,0) + (1,0,2) = (3,0,2)$$

$$\mathbf{Need}_1 = (1,2,2) - (1,0,2) = (0,2,0)$$

- So, we arrive at the following new state

Is the system in safe state?

Yes. Because we find a safe sequence: $\langle P_1, P_3, P_4, P_0, P_2 \rangle$.

Grant the request of P_1 .

What about request of (3 3 0) resources for P_4 ?

What about request of (0 2 0) resources for P_0 ?

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 2	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	



IT'S OVER

IT'S DONE